

Programming Assignment 1: REPORT

Optimization of Double Precision Floating Point Matrix Multiply Accumulator

- David Garcia, Rahul Godbole

The goal of this project was for us to optimize the matrix multiply accumulate operation $C = C + A*B$. All matrices are 64-bit double precision floating type. The computation is performed on a machine with 32kB L1 cache and 4MB L2 cache.

Following optimizations were tried out:

1. **Blocking (+)**: Chunks of sub matrices computed with varying block sizes. Specialized matrix sizes used multiple blocking levels. Optimum Block Size for final implementation found to be blocks of 256 for the outermost loop, 16 for the middle loop. This was mostly a trial and error approach. The goal was for us to exploit spatial locality in the L1 cache based on the size of the blocks.
2. **Prefetching (-)**: Did not seem to help us much. Compiler was already implementing some prefetching. Also, the HW prefetcher was doing its job.
3. **Transpose (+)**: Allows us to access memory sequentially i.e. data in the same cache line.
4. **SSE2 intrinsics (+)**: Maximum benefit seen with SIMD extension intrinsics. We used the `_mm_load_pd`, `_mm_add_pd`, `_mm_mul_pd` and `_mm_set1_pd` intrinsic functions.
5. **Compiler flags (+)**: Several of these were used in the project.
 - (a) `-funroll-loops`: Looking at the Intel compiler reference guide, this option is OFF by default. In conjunction with `-march=corei7`, there was a slight degradation in performance - this is also acknowledged in the Intel compiler reference guide.
 - (b) `-march=corei7`: This option enables SSE, SSE2, SSE3 intrinsics. Hence, is required for the SSE2 optimizations to take effect.
 - (c) `-std=gnu99`: no effect, was hoping it would better enforce `__restrict__`
 - (d) `-opt-args-in-reg=all`: Not much of an impact. Causes function parameters to be passed in registers (instead of on the stack), where the function may or may not be in the compilation unit scope.
 - (e) `-fprefetch-loop-arrays`: increased the N=128 case to 4500 MFLOPS, decreased the rest
 - (f) `--param l1-cache-line-size, l1-cache-size, l2-cache-size`: no effect

6. Strassen Algorithm: This algorithm reduces the number of multiplications in exchange for additions-- for block sizes 1024, 2048, where improvement would have been seen, it resulted in slowdown. (See code in appendix A). It was necessary to use our algorithm from before changing order of operations (see #7) as the new algorithm was less efficient for multiplying matrices without summing to the result matrix.

7. Changing Order of operations: Obtained speedup by iterating through C and B multiple

times while keeping A's values stored every iteration of the middle loop. This rendered transposition of B unnecessary.

8. `__restrict` keyword: Added this to pointer declarations to allow compiler optimizations that rely on multiple pointers referencing overlapping data, but it had no observable effect.


9. Moving operations as far out of loops as possible: This was done in case the compiler or prefetcher couldn't optimize for strides of $i*N$ --instead the array pointers for eg $A[i*n+k]$ were moved to $\text{double* } a = A + i*n$, and then referenced as $a[k]$. This gave no observable speedup, so the compiler optimizations were probably working.

Results

Shown below are the performance numbers we observed on running our optimized program.

Dim.	MFLOPS	Runtime	0x80000032	0x8000003b	0x80000000	0x80000007
2	1556.507	0.0000	49	24	0	0
4	1326.273	0.0001	505	224	0	0
8	2731.304	0.0004	2109	871	0	0
16	3708.236	0.0022	12645	5130	0	0
32	4004.722	0.0164	80368	38010	27	0
64	4081.402	0.1285	562659	299508	2962	0
128	4122.454	1.0174	4183379	2374191	21592	0
256	4060.898	8.2628	32199215	19120767	1310396	293
512	3684.181	72.8616	257592982	169475282	17715403	137782
1024	3639.575	590.0370	2060742586	1367032028	141584878	1364390
2048	3661.295	4692.2930	16485936910	10892166149	1133779672	11282069

1. Graph



2. Table

Dimension	MFLOPS OPTIMIZED	MFLOPS Naïve
2	1556.501	765.024
4	1326.273	763.071
8	2731.304	769.536
16	3708.236	760.198
32	4004.722	684.072
64	4081.402	575.108
128	4122.454	331.329
256	4060.898	227.584
512	3684.181	120.259
1024	3639.575	97.874
2048	3661.295	

The numbers shown in the above table were observed on average after multiple runs, hence they are representative of the mean.

Appendix A:

```

double *R = malloc(N*N*sizeof(double));
double *M1 = malloc(n*n*sizeof(double));
double *M2 = malloc(n*n*sizeof(double));
double *M3 = malloc(n*n*sizeof(double));
double *M4 = malloc(n*n*sizeof(double));
double *M5 = malloc(n*n*sizeof(double));
double *M6 = malloc(n*n*sizeof(double));
double *M7 = malloc(n*n*sizeof(double));
double *tmp1 = malloc(n*n*sizeof(double));
double *tmp2 = malloc(n*n*sizeof(double));

const double *A12 = A + n, *A21 = A+N*n, *A22 = A + (N+1)*n;
const double *B12 = Bp + n, *B21 = Bp+N*n, *B22 = Bp + (N+1)*n;
double *R12 = R + n, *R21 = R+N*n, *R22 = R + (N+1)*n;

addM(n, A, A22, tmp1, N, N, n);
addM(n, Bp, B22, tmp2, N, N, n);
mulM(n, tmp1, tmp2, M1, n, n, n);
addM(n, A21, A22, tmp1, N, N, n);
mulM(n, tmp1, Bp, M2, n, N, n);
subM(n, B12, B22, tmp1, N, N, n);
mulM(n, A, tmp1, M3, N, n, n);
subM(n, B21, Bp, tmp1, N, N, n);
mulM(n, A22, tmp1, M4, N, n, n);

```

```

    addM(n, A, A12, tmp1, N, N, n);
    mulM(n, tmp1, B22, M5, n, N, n);
    subM(n, A21, A, tmp1, N, N, n);
    addM(n, Bp, B12, tmp2, N, N, n);
    mulM(n, tmp1, tmp2, M6, n, n, n);
    subM(n, A12, A22, tmp1, N, N, n);
    addM(n, B21, B22, tmp2, N, N, n);
    mulM(n, tmp1, tmp2, M7, n, n, n);

    addM(n, M1, M4, tmp1, n, n, n);
    subM(n, tmp1, M5, tmp2, n, n, n);
    addM(n, tmp2, M7, R, n, n, N);
    addM(n, M3, M5, R12, n, n, N);
    addM(n, M2, M4, R21, n, n, N);
    subM(n, M1, M2, tmp1, n, n, n);
    addM(n, M3, M6, tmp2, n, n, n);
    addM(n, tmp1, tmp2, R22, n, n, N);
    addMTo(N, R, C, N, N);
}
}

void addM(int N, const double*__restrict__ A, const
double*__restrict__ B, double *__restrict__ C, int incA, int incB,
int incC) {
    int i, j;
    int n = N/2;
    int inA = incA/2;
    int inB = incB/2;
    int inC = incC/2;
    __m128d* a = (__m128d*) A;
    __m128d* b = (__m128d*) B;
    __m128d* c = (__m128d*) C;
    for (i = 0; i < N; i++)
        for (j = 0; j < n; j++)
            c[i*inC + j] = _mm_add_pd(a[i*inA + j], b[i*inB + j]);
}

void addMTo(int N, const double*__restrict__ A, double *__restrict__
C, int incA, int incC) {
    int i, j;
    int n = N/2;
    int inA = incA/2;
    int inC = incC/2;

```

```

    __m128d* a = (__m128d*) A;
    __m128d* c = (__m128d*) C;
    for (i = 0; i < N; i++)
        for (j = 0; j < n; j++)
            c[i*inC + j] += a[i*inA + j];
}

void subM(int N, const double*__restrict__ A, const
double*__restrict__ B, double *__restrict__ C, int incA, int incB,
int incC) {
    int i, j;
    int n = N/2;
    int inA = incA/2;
    int inB = incB/2;
    int inC = incC/2;
    __m128d* a = (__m128d*) A;
    __m128d* b = (__m128d*) B;
    __m128d* c = (__m128d*) C;
    for (i = 0; i < N; i++)
        for (j = 0; j < n; j++)
            c[i*inC + j] = _mm_sub_pd(a[i*inA + j], b[i*inB + j]);
}

void mulM (int N, const double*__restrict__ A, const
double*__restrict__ Bp, double*__restrict__ C, int incA, int incB,
int incC) {
    int i, j, k, ii, jj, kk, iii, jjj, kkk;
    int n = N/2;
    int inA = incA/2;
    int inC = incC/2;
    __m128d* mA = (__m128d*) A;
    __m128d* mC = (__m128d*) C;

    union pd_d {
        __m128d pd;
        double d[2];
    } r1;
    double *B = malloc(N*N*sizeof(double));
    __m128d* mB = (__m128d*) B;
    // transpose array to line up memory access
    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            B[i*N+j] = Bp[j*incB+i];

```

```

int inB = N/2;

// for (jj = 0; jj < N; jj += BSJ)
  for (i = 0; i < N; ++i)
    for (j = 0; j < N; ++j) {
//      for (j = jj; j < min(jj+BS,N); ++j) {
        // sse2 register to accumulate 2 additions at a time,
        // for every 2*kth and 2*k+1th array element (= kth element
        // of __m128d* array)
        r1.pd = _mm_set_pd(0.0, 0.0);
        for (k = 0; k < n; ++k) {
            // multiply elements from mA and mB and add to cumulative
total
                                r1.pd = _mm_add_pd(    r1.pd,
_mm_mul_pd(mA[i*inA+k],mB[j*inB+k]));
        }
        C[i*incC + j] = r1.d[0] + r1.d[1];
    }

}

```

Output:

Dim.	MFLOPS	Runtime	0x80000032	0x8000003b	0x80000000	0x80000007
2	1612.141	0.0000	50	23	0	0
4	1155.054	0.0001	507	259	0	0
8	2102.033	0.0005	2111	1136	0	0
16	3701.522	0.0022	12647	5140	0	0
32	4000.786	0.0164	80368	38053	22	0
64	3701.244	0.1417	562651	328972	2975	0
128	4097.581	1.0236	4183367	2379841	21540	0
256	4077.076	8.2300	32199187	19180943	1307174	131
512	3698.681	72.5760	257592731	169259916	17675235	131711
1024	2925.544	734.0460	1669939669	1714349160	121877518	6042034
2048	987.968	17389.0860	13020613551	40373901775	960540039	818772448

qsub: job 2022296.cyclades-master.stanford.edu completed